

# Benchmarking Horovod and Ray for Distributed DNN training on GPU cluster

Shreeparna Dey    Roopkatha Banerjee

April 28, 2022

# Distributed Deep Learning

# Need for Distributed Learning?

- Data set too big.
- Model size too big.
- Results need to be produced in a reasonable amount of time.

# Methods of Distributed Learning

- Model Parallelism - difficult to split model efficiently.
- Data Parallelism - limited by the all-reduce at the end of each iteration

# Distributed Deep Learning Frameworks

# Horovod

- Optimizes the inter-node communication → Data Parallelism
- Uses NCCL 2 → Provides semantics for inter-node ring-allreduce.
- No fault tolerance

# Ray Train

- Facilitates inter-node communication with the use of a parameter server
- Parameter server collects and averages the gradients
- Allows fault tolerance

# Experiments



# Datasets

We have used three sets of data: MNIST, FMNIST and Data from the Sloan Digital Sky Survey.

## **MNIST Dataset**

- image dataset of handwritten digits from 0 to 9, multiclass-classification problem with 10 class labels.
- 60000 training and 10000 testing data points
- one image which is in grayscale is 28x28

## **Fashion-MNIST Dataset**

- image dataset of Zalando's article images, multiclass-classification problem with 10 class labels.
- 60000 training and 10000 testing data points
- one image which is in grayscale is 28x28

## **Sloan-Digital-Sky-Survey Dataset**

- imaging and spectroscopic red-shift survey of the northern and southern hemispheres.

# Models

- **Resnet50:** It is a DNN (Deep Neural Network) which is 50 layers deep where there is 48 Convolution layers and 1 Max- Pooling and 1 Average Pooling layer
- **MobileNetV2:** MobileNet-v2 by Google is a convolutional neural network that is 53 layers deep. It is a light-weight feature detector, which is suited to devices with low computational power

# Hardware Setup

- Node 5 of the IOE cluster which houses two Titan RTXs with a total of 48GB GPU memory and two AMD EPYC processors with a total 32 threads and 256GB RAM
- Node 9 of the IOE cluster which house four V100 with a total of 128GB memory and an Intel Xeon processor with 32 threads and 192GB RAM. These cores have Non-volatile memory and have  $2 \times 128\text{GB}$  Intel Optane memory installed.

Each of the nodes on the IOE cluster is running CentOS7, gcc version 9.3 and CUDA version 11.0. The cluster uses PBS job scheduler.

# Results

# Training Time

batch size = 32	RayTrain		Training time(Sec)	
	Node 5(2*TitanRTX)		Node 9(4*V100)	
	1	2	1	2
RN50+MNIST	4.0798	2.789	3.353	1.848
MN+F-MNIST	0.817	0.654	0.792	0.808
batch size = 32	Horovod		Training time(Sec)	
	Node 5(2*TitanRTX)		Node 9(4*V100)	
	1	2	1	2
RN50+MNIST	4.132	3.4474	3	3.7002
MN+F-MNIST	0.691	0.6409	0.614	0.936

**Figure:** Model Training Times of models on different datasets across RayTrain and Horovod frameworks with different number and variety of GPUs

# Accuracy

batch size = 32	RayTrain		Testing Accuracy	
	Node 5(2*TitanRTX)		Node 9(4*V100)	
	1	2	1	2
RN50+MNIST	0.9883	0.9922	0.99	0.99
MN+F-MNIST	0.6951	0.47	0.5802	0.7498
batch size = 32	Horovod		Testing Accuracy	
	Node 5(2*TitanRTX)		Node 9(4*V100)	
	1	2	1	2
RN50+MNIST	90.2	90.4	90.3	90.3
MN+F-MNIST	72.9	72	72.2	71.8

**Figure:** Accuracies of models on different testing datasets across RayTrain and Horovod frameworks with different number and variety of GPUs

# Images processed per Second

batch size = 32	RayTrain		Images/Sec	
	Node 5(2*TitanRTX)		Node 9(4*V100)	
	1	2	1	2
RN50+MNIST	14706	21513	17894	32467
MN+F-MNIST	73739	91743	75757	74257
batch size = 32	Horovod		Images/Sec	
	Node 5(2*TitanRTX)		Node 9(4*V100)	
	1	2	1	2
RN50+MNIST	14520	17404	17400	16215
MN+F-MNIST	86830	93618	93618	64102

**Figure:** Images processed per Second by models on different datasets across RayTrain and Horovod frameworks with different number and variety of GPUs

# Conclusions, Challenges. and Future Work

- In our experiments, Ray Train shows good speedups for multi-GPU training, especially when compared to Horovod.
- Caffe2 has been deprecated and merged into PyTorch
- We also tried a PyTorch distributed implementation, however, we could not debug it in time.
- We also tried to train big models like VGG16( $\sim 138$  billion parameters) but it could not fit into the available GPU memory.



# Thank you!